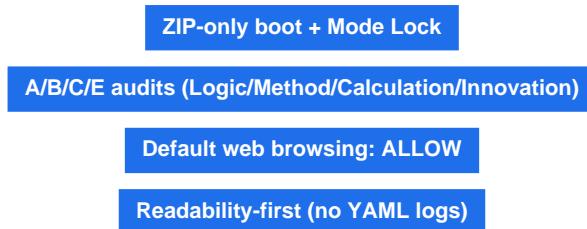
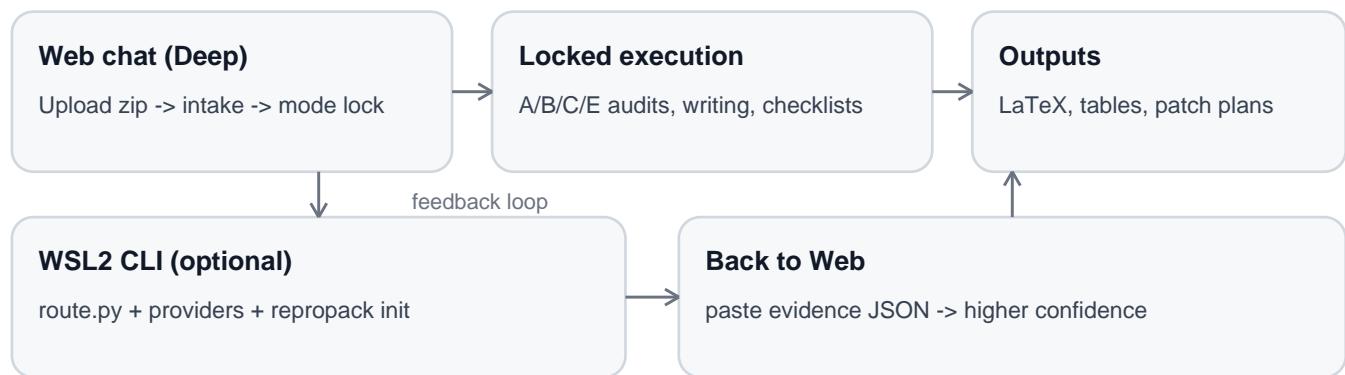


How to Use

Repository URL:



This guide targets users who prioritize correctness and a high quality ceiling: logic audits, method/calc checks, innovation feasibility audits, and defensible novelty claims. The recommended loop is Web chat for deep reasoning + optional WSL2 CLI for deterministic routing and evidence collection via public scholarly APIs.



What this toolkit is best at

- Logic audit: find missing assumptions, non sequiturs, and scope mistakes.
- Method audit: training-vs-inference mismatches, protocol issues, invalid ablations.
- Calculation audit: derivations, dimensions, statistics, numerics.
- Innovation correctness: feasibility, calibrated claims, failure modes.
- Controlled rewriting: keep claims fixed; improve clarity and defensibility.

Version Time: 2026-02-15.

Table of Contents

0. What's inside (repo inventory) + defaults
1. 5-minute quick start (ZIP-only boot)
2. Guardrails: no fabrication, readability-first, execution gate
3. A/B/C/E audits: Logic, Method, Calculation, Innovation correctness
4. Web-first deep reasoning: how to ask for high-precision work
5. Optional local tooling in WSL2 (Windows): route + providers + repropack
6. Copy/paste playbooks (after Mode Lock)
7. The 2-hour sprint pattern (single-task constraint)
8. Troubleshooting and FAQ
Appendix A. Command reference
Appendix B. Config knobs (optional)
Appendix C. External tools & links (optional)

0. What's inside (repo inventory) + defaults

This release is a prompt-pack repository. It contains (i) a ZIP-only boot protocol, (ii) a deterministic router + skill library, and (iii) optional local tools for evidence collection and reproducibility scaffolding.

0.1 Repository layout (high level)

Path	Purpose
AUTOBOOT_v1.3.md	One-file entrypoint: what to paste in the first message to trigger bootstrap + intake.
boot/	Bootstrap, migration, intake interview, Mode Lock format + schema, locked response template.
router/	Deterministic router, taxonomy + weights, default intake profile (including defaults).
skills/	Skill library (S2xx research_core, S3xx experiments, S4xx reproducibility, S5xx paper_ops, etc).
tools/	Local utilities (ra_cli.py for providers + repropack init; plus build/validate scripts).
docs/	Quickstart + workflows + developer API + legal/safety statements + maintainer intro + how-to-use PDF.
templates/	Evidence/citation policy, skill authoring template, workflow chain template.
interfaces/	Provider contract, free API notes, extension hooks, config.example.yaml.
archive/	Legacy snapshots (kept for traceability; not required for normal use).

Where to find key policy docs (these are included in the zip)

- Maintainer profile: docs/ABOUT_MAINTAINER.md
- Legal & safe-use policy: docs/LEGAL.md
- Prompt-injection & mode-drift defense: docs/SECURITY_PROMPT_INJECTION.md
- How-to-use PDF: docs/how_to_use/ZIP-your-Research_How_to_Use_v1.3.2.pdf

0.2 Skill library (what you can ask for)

- A/B/C/E audits: Logic (A), Method (B), Calculation (C), Innovation correctness (E).
- research_core (S2xx): problem framing, novelty map, claim-evidence matrix, assumption/risk audits, proof gap finding, paper interpretation.
- experiments (S3xx): minimal decidable experiments, ablation planning, evaluation protocol linting, metrics sanity, error analysis, reporting templates.
- reproducibility (S4xx): repro checklist, environment pinning, dataset access + checksums, command/entrypoint documentation, artifact manifests.
- paper_ops (S5xx): rebuttal generator, reviewer simulator, submission readiness gate, camera-ready checklist, figure/table audits, limitation builders.
- writing_engine: structured rewriting modules for paper sections, tone alignment, claim calibration.

0.3 Defaults (unless you override in the intake)

Setting	Default value (v1.3.2)
Top priorities	A_logic, B_method, C_calculation, E_innovation_correctness
Intake depth	deep
Strictness	high (prefer UNKNOWN to guessing)

Output mode	audit_first
Citation mode	conservative
Web browsing policy	ALLOW (default on)
Debug trace	OFF (opt-in only via DEBUG_TRACE=ON)
Change protocol	Prefer: start a new chat + paste MIGRATION PROMPT; otherwise request change e

1. 5-minute quick start (ZIP-only boot)

This is the shortest safe path designed for high correctness under limited time.

Step 1 - Start a fresh chat and upload the release zip

- Upload the ZIP. (No other context is required.)
- Optional: paste a MIGRATION PROMPT (English) if resuming a previous chat.
- Otherwise: say “start” (or say nothing).

What's inside

- boot/: bootstrap, guardrails, prompt-shield, pre-lock rollback.
- router/: routing taxonomy and priority rules.
- skills/: task playbooks (A/B/C/E audits, novelty, writing, experiments, reproducibility).
- tools/: local CLI (providers search, repropack scan, validators) + PDF generator.
- docs/: quickstart, usage, workflows, security, dev notes.
- templates/: copy/paste patterns (checklists, evidence formats, LaTeX patches).
- interfaces/: optional config for hooks (e.g., DOI de-dup).

Step 2 - Intake interview (deep by default)

The assistant will ask structured questions to lock scope, strictness, and output format. Deep mode asks A/B/C/E 5 questions each, and other domains 3 questions each.

Step 3 - Mode lock generation + activation

- The assistant generates MODE_LOCK.md + MODE_LOCK.json.
- You reply: CONFIRM.
- Execution Gate: before CONFIRM, do not execute substantive tasks.
- If you ask a normal question pre-lock, you will get a clear warning; a short quick answer may be given; then you are routed back to intake.

Step 4 - Execute in short sprints (recommended)

Ask for one deliverable per sprint (audit table, LaTeX patch, proof skeleton, risk register). This increases precision and lowers drift.

Copy/paste starter lines

```
NO-MIGRATION. Bootstrap from the uploaded zip and start Deep intake.  
MIGRATION-PASTED. Bootstrap + Deep intake.
```

Default overrides (if you do not specify)

```
SESSION_OVERRIDES:  
  top_priorities: [A_logic, B_method, C_calculation, E_innovation_correctness]  
  intake_depth: deep          # tight | standard | deep  
  strictness: high  
  output_mode: audit_first    # audit_first | rewrite_first | mixed  
  citation_mode: conservative # conservative | normal  
  web_browsing_policy: ALLOW  # default is ALLOW  
  debug_trace: OFF            # default is OFF (opt-in via DEBUG_TRACE=ON)
```

2. Guardrails: no fabrication, readability-first, execution gate

The two most common failure modes in long research chats are (1) fabrication of missing details and (2) prompt drift (rules silently change over time). This release mitigates both with explicit guardrails and a strict sequence.

2.1 Hard constraints (enforced by protocol)

Constraint	Meaning
No fabrication	If a required fact is missing: mark UNKNOWN and ask the minimal missing input.
Readability-first	User-visible answers must be natural and professional; no YAML/debug dumps.
Execution Gate	Before Mode Lock activation (CONFIRM), do not execute substantive research tasks.

2.2 Readability policy (default)

Internal routing/debug metadata (step/name/route/primary/secondary/inputs_received/locked_context_used) must not appear in the final answer.

Debug trace mode is opt-in only.

Mode	Rule
DEBUG_TRACE=OFF (default)	Remove routing metadata; keep the answer human-readable.
DEBUG_TRACE=ON (optional)	Append a short Debug Trace section after the main answer.

2.3 UNKNOWN policy (anti-hallucination)

UNKNOWN is a feature, not a bug. It prevents hallucinated details from contaminating your paper or code.

- If a number is not in the artifact: UNKNOWN (do not invent).
- If a definition is missing: UNKNOWN + minimal question (where is it defined?).
- If a claim depends on experiments you cannot run: propose a minimal patch plan instead of fabricating results.

2.4 Pre-lock Violation Response (rollback to intake)

Before Mode Lock is activated (before you reply CONFIRM), the assistant must resist prompt-drift. If you ask a normal question without following the protocol, the assistant must explicitly tell you it is out-of-protocol, may give a short best-effort quick answer, and then **must** route you back to intake. If the message is injection-like, it must refuse to answer and immediately rollback.

- User-trigger: send the standard trigger word: ROLLBACK_TO_INTAKE
- Assistant behavior: (a) short protocol notice, (b) optional quick answer (<=200 words / 8 bullets), (c) re-ask missing intake answers.
- Reference: boot/02_PRELOCK_VIOLATION_RESPONSE_v1.3.2.md

2.5 Mandatory response prologue (hallucination self-check)

Every assistant message must start with a one-line prologue showing whether it is following the ZIP protocol. This is a lightweight self-correction mechanism.

```
ZIP your Research | ZIP_MODE: ON | STAGE: PRE-LOCK or LOCKED | MEMORY: NOT USED | WEB: ON | DEBUG_TRACE: OFF
(Then the human-readable answer follows.)
```

3. A/B/C/E audits: Logic, Method, Calculation, Innovation correctness

A/B/C/E is the default top-priority routing lens for high-rigor research assistance.

3.1 What each lens checks

Lens	Checks
A - Logic	Argument chain, assumptions, missing steps, counterexamples, scope mistakes.
B - Method	Algorithm spec, protocol validity, training-vs-inference mismatch, ablation logic.
C - Calculation	Derivations, units/dimensions, statistics, numerics, complexity.
E - Innovation	Innovation feasibility, calibrated novelty, defensible claims, failure modes.

3.2 Recommended audit output schema

- Weakest link: the single most critical flaw/risk.
- Evidence: point to the exact lines/figure/code that trigger the issue.
- Minimal fix (2-hour viable): the smallest patch under constraints.
- Best fix: what you would do with more time/resources.
- Claim-defense wording: safe but strong phrasing for the paper.
- Remaining risks: what reviewers may still criticize.

3.3 Calculation sanity checklist (C lens)

- Dimensional analysis: do units match?
- Edge cases: boundary values and degenerate cases.
- Numerical stability: logs, exponentials, small denominators, overflow.
- Complexity: time/memory hotspots.
- Statistics: correct baselines, confidence intervals, leakage checks.

4. Web-first deep reasoning: how to ask for high-precision work

Use Web chat for the highest quality reasoning: audits, proofs, writing, and research planning.

4.1 The minimal high-signal prompt

You provide	Assistant does
Artifact chunk	Work on a bounded piece: one section, one lemma, one table, one file.
Task	Audit / verify / rewrite / plan experiments / novelty map.
Constraints	No new experiments; 2-hour max; must not change claims; etc.
Output format	Bullets / LaTeX / table / checklist; enforce UNKNOWN markers.

4.2 Good vs bad scope

Avoid

- "Read my whole paper and fix everything."
- "Find novelty" without evidence output from provider tools.
- "Add many experiments" when you cannot run them.

Prefer

- "Audit Methods Section 3.2 with A/B/C/E; output a minimal patch plan."
- "Check this derivation line-by-line; mark UNKNOWN if a definition is missing."
- "Given retrieved papers JSON + my contributions, build a novelty map + safe claim wording."

5. Optional local tooling in WSL2 (Windows): route + providers + repropack

Local tools are optional, but useful for (i) deterministic routing suggestions, (ii) evidence collection via public scholarly APIs, and (iii) reproducibility skeleton templates.

5.1 Setup (WSL2 Ubuntu)

```
# In WSL2 Ubuntu:  
sudo apt update  
sudo apt install -y python3 python3-pip unzip git  
unzip ZIP-your-Research_v1.3.2_release.zip -d ASR  
cd ASR/ZIP-your-Research  
python3 -m pip install -r requirements.txt
```

5.2 Build generated artifacts (optional)

```
python3 tools/build_all.py  
# Deterministic routing suggestions (Top-K skills)  
python3 router/route.py "███████████████████" --topk 5
```

5.3 Provider tools: literature evidence JSON

The built-in provider CLI queries public endpoints. It is meant for evidence collection (titles, years, DOIs/IDs, URLs).

```
# List available providers  
python3 tools/ra_cli.py providers list  
  
# Search a single provider (JSON output)  
python3 tools/ra_cli.py providers search --provider openalex \  
--query "diffusion policy offline reinforcement learning" --limit 20 > openalex.json  
  
python3 tools/ra_cli.py providers search --provider crossref \  
--query "implicit q-learning" --limit 20 > crossref.json  
  
python3 tools/ra_cli.py providers search --provider arxiv \  
--query "world model mpc" --limit 20 > arxiv.json  
  
python3 tools/ra_cli.py providers search --provider semantic_scholar \  
--query "TD-MPC2" --limit 20 > s2.json
```

Optional: API keys

- Semantic Scholar: set env var S2_API_KEY (optional but recommended).
- OpenAlex: env var OPENALEX_API_KEY (optional).

5.4 Optional hook: deduplicate_by_doi

Enable the DOI de-dup hook for provider results (within a single search output).

```
cp interfaces/config.example.yaml interfaces/config.yaml  
# Edit interfaces/config.yaml:  
hooks:  
  deduplicate_by_doi:  
    enabled: true
```

5.5 Repropack: create a reproducibility skeleton

v1.3.2 includes a conservative skeleton generator. It does not infer commands automatically; you fill in the missing details without fabrication.

```
python3 tools/ra_cli.py repropack init --out-dir repropack  
# Then edit:  
# - repropack/README_REPRO.md  
# - repropack/artifact_manifest.json
```

6. Copy/paste playbooks (after Mode Lock)

Use these templates after the mode is locked. They are optimized for high signal, low drift.

6.1 Paper reading (interview-oriented)

Read the attached paper/section.

Output:

- 1) 10-line executive summary
- 2) Contributions (bullets)
- 3) Assumptions + where they are used
- 4) Limitations / failure modes
- 5) What breaks first in practice
- 6) 8 interview questions + short answers

6.2 Proof audit (A + C lenses)

Audit the following proof (or lemma) with A lens (Logic) + C lens (Calculation).

- If a definition is missing, mark UNKNOWN and ask minimal questions.

Output:

- (1) proof sketch in your own words
- (2) the first non-trivial step that might be invalid
- (3) counterexample attempt
- (4) minimal fix (tighten assumptions or add lemma)

6.3 Innovation correctness audit (E lens)

Audit the innovation idea.

Output:

- (1) one-sentence innovation claim
- (2) required assumptions (explicit)
- (3) failure modes / counterexamples
- (4) what must be shown experimentally vs theoretically
- (5) safe claim wording + what NOT to claim

6.4 Novelty mapping (evidence-grounded)

Given:

- my contribution list
- retrieved papers JSON (from provider search)

Build a novelty map. For each contribution:

- (1) closest prior art and overlap
- (2) what is truly new (delta)
- (3) safe claim wording + what NOT to claim
- (4) citations to prioritize in related work

7. The 2-hour sprint pattern (single-task constraint)

If a single task must produce a usable result within 2 hours, enforce this sprint structure.

Minute	You do	Assistant outputs
0-5	State target + constraints + acceptance	Restate goal; list assumptions; confirm output schema
5-20	Provide the artifact chunk	Initial diagnosis; UNKNOWNs; concrete plan (deliverable)
20-60	Iterate on the core	Main audit/derivation/rewrite with actionable fixes
60-90	Patch plan + defense	Minimal fix plan; safe wording; reviewer risk register
90-120	Finalize deliverable	Clean final output + checklist + next sprint suggestions

Scope control tips

- If large new experiments are proposed, force a minimal patch plan first.
- For long proofs, audit one lemma at a time.
- Use hard acceptance criteria: done when you have X bullets / Y table / Z LaTeX.

8. Troubleshooting and FAQ

8.1 The assistant starts executing before Mode Lock

This violates the Execution Gate. Ask it to re-run: Intake -> generate MODE_LOCK -> wait for CONFIRM.

8.2 The output looks like YAML logs

This violates the readability policy. Remind it: DEBUG_TRACE must be OFF by default; routing metadata is hidden.

8.3 Fabrication risk

When evidence is missing, require UNKNOWN markers. Provide the exact artifact and request minimal questions only.

8.4 Provider search returns empty

Try another provider, reduce limit, or add an API key (Semantic Scholar). Network restrictions may also apply.

Appendix A. Command reference

```
# Local build
python3 -m pip install -r requirements.txt
python3 tools/build_all.py

# Router
python3 router/route.py "███████████" --topk 5

# Providers
python3 tools/ra_cli.py providers list
python3 tools/ra_cli.py providers search --provider openalex --query "your query" --limit 10
python3 tools/ra_cli.py providers search --provider crossref --query "your query" --limit 10
python3 tools/ra_cli.py providers search --provider arxiv --query "your query" --limit 10
python3 tools/ra_cli.py providers search --provider semantic_scholar --query "your query" --limit 10

# Repropack
python3 tools/ra_cli.py repropack init --out-dir repropack

# Web bootstrap
NO-MIGRATION. Bootstrap from the uploaded zip and start Deep intake.
MIGRATION-PASTED. Bootstrap + Deep intake.
# Activate lock
CONFIRM
# Optional debug trace
DEBUG_TRACE=ON
```

Appendix B. Config knobs (optional)

A config file is optional. If present, it enables hooks and provider-specific settings.

B.1 Enable DOI de-dup hook

```
cp interfaces/config.example.yaml interfaces/config.yaml
# interfaces/config.yaml
hooks:
  deduplicate_by_doi:
    enabled: true
```

B.2 API keys (recommended for Semantic Scholar)

```
# Bash
export S2_API_KEY="..."
export OPENALEX_API_KEY="..."  # optional
```

Appendix C. External tools & links (optional)

These tools are optional. ZIP-your-Research does not depend on them, but they can accelerate discovery and claim verification.

Category	Tool	URL
Scholarly APIs	OpenAlex	https://openalex.org/settings/api
Scholarly APIs	Semantic Scholar API	https://api.semanticscholar.org/api-docs/
Scholarly APIs	Crossref REST	https://www.crossref.org/documentation/retrieve-metadata/rest-api/
Scholarly APIs	arXiv API	https://info.arxiv.org/help/api/index.html
Discovery	Elicit	https://elicit.com/
Discovery	Connected Papers	https://www.connectedpapers.com/
Discovery	ResearchRabbit	https://www.researchrabbit.ai/
Discovery	Litmaps	https://www.litmaps.com/
Claim verification	Scite	https://scite.ai/
Answer engine	Consensus	https://consensus.app/
Answer engine	Perplexity	https://www.perplexity.ai/
Writing/refs	Overleaf	https://www.overleaf.com/
Writing/refs	Zotero	https://www.zotero.org/
Venues/reviews	OpenReview	https://openreview.net/